

KEDAR LIMBALKAR

Data Scientist | ML Engineer

kedarlimbalkar@gmail.com | +91-9284374129 | Pune, Maharashtra | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Machine Learning Engineer and Data Scientist with hands-on expertise in deep learning, NLP, LLMs, computer vision, and end-to-end ML pipelines. Proficient in RAG workflows, model fine-tuning, prompt engineering, MLOps, and Docker deployment. Proven track record of building and deploying production-grade AI systems with measurable accuracy and latency improvements.

TECHNICAL SKILLS

Languages & Libraries: Python (Pandas, NumPy, Scikit-learn, XGBoost, Keras, PySpark, Matplotlib, Seaborn), SQL

ML / Deep Learning: TensorFlow, CNN (Convolutional Neural Networks), Computer Vision, Feature Engineering, Statistical Analysis

Generative AI & LLMs: LangChain, LangGraph, FAISS, Llama 3.3, Hugging Face, Prompt Engineering, Model Fine-Tuning, RAG (Retrieval-Augmented Generation), Vector Embeddings, Transformer Models, Semantic Search, Multi-Agent Systems

NLP: Natural Language Processing, Text Classification, Sentiment Analysis, Named Entity Recognition, Tokenization, Text Preprocessing, Word Embeddings

MLOps & DevOps: Docker, Git, GitHub, CI/CD, MLOps Deployment, Streamlit, OpenCV, REST API, Apache Airflow

Data Engineering: ETL Pipelines, Data Modeling, Data Wrangling, Data Visualization, EDA (Exploratory Data Analysis)

Databases & Cloud: MySQL, BigQuery, Snowflake

WORK EXPERIENCE

Intucate | Business Analyst Intern

January 2026 – April 2026

- Analyzed content asset performance metrics using data analysis and EDA techniques, identifying trends that improved educational engagement outcomes by 15% for target demographics.
- Collaborated with product and engineering teams to translate business requirements into structured workflows, aligning roadmap priorities across 3+ departments.

PROJECTS

Conversational AI for Retail | [GitHub](#) | [Drive](#)

January 2026 – February 2026

Technologies: Python, LangGraph, LangChain, Llama 3.3, FAISS, RAG, Vector Embeddings, NLP, Prompt Engineering, Multi-Agent Systems

- Preprocessed 44,884 retail records using a 3,000/500 token chunking strategy to build a FAISS vector knowledge base, achieving 100% semantic integrity for rapid retrieval.
- Designed a multi-agent LangGraph RAG pipeline achieving 92% domain accuracy and sub-1.2s query latency, and applied prompt engineering and fine-tuning with Llama 3.3 to reduce hallucination rate by 18% across 10,000+ interactions.

Emotion Recognition System | [GitHub](#) | [Drive](#)

February 2026 – March 2026

Technologies: Python, Keras, TensorFlow, CNN, OpenCV, Streamlit, Computer Vision, Deep Learning, Transfer Learning, Batch Normalization

- Built a CNN in Keras to classify 7 emotional states from 35,000+ facial images, achieving 88% test accuracy using Dropout and Batch Normalization layers, and improving precision by 12% over the baseline architecture.
- Deployed a real-time computer vision inference pipeline with OpenCV and Streamlit at 30+ FPS, applying transfer learning and hyperparameter tuning to reduce overfitting by 20%.

Shopzilla CSAT Prediction | [GitHub](#)

April 2026 – May 2026

Technologies: Python, XGBoost, Scikit-learn, Pandas, NumPy, Matplotlib, Seaborn, Streamlit, RandomizedSearchCV

- Built an end-to-end CSAT prediction pipeline over 85,900+ e-commerce support tickets, engineering response-time and remark-length features and tuning an XGBoost classifier via RandomizedSearchCV with stratified 5-fold cross-validation.
- Optimized for weighted F1 (~0.61) over raw accuracy to address severe class imbalance (69% majority class), and deployed the trained model through an interactive Streamlit app for real-time ticket scoring.

EDUCATION

MS in Computer Science: Machine Learning & AI Engineering — Woolf University

2025 – 2027

Relevant Coursework: Machine Learning, Artificial Intelligence, Deep Learning, NLP, Data Science, Statistical Modeling

B.Tech in Computer Engineering — Government College of Engineering, Jalgaon

2021 – 2025

CERTIFICATIONS & ACHIEVEMENTS

Full Stack Data Science and AI | AlmaBetter | [View Certificate](#)

- Completed 6 production-ready ML/DL projects across 8+ datasets spanning machine learning, deep learning, NLP, and Docker-based MLOps deployment.
- District Winner — National Level Computing Intelligence Olympiad (CIO): ranked 1st in logical reasoning and data interpretation among 500+ participants.
- Core Member, Technical Council 2025 — coordinated 3 national-level events, managing logistics for 200+ attendees.